

Count-Only Inverse-Variance Weighting for Intervention Harvesting

Anonymous Author(s)

Abstract

Intervention-harvesting estimators of position bias [1] combine per-cell click-rate estimates across many $(q, d, \text{position})$ cells whose impression counts span several orders of magnitude. Under the Position-Based Model (PBM), with cell-level relevance treated as homoscedastic, the *harmonic* weight $\omega_i = N_k^i N_{k'}^i / (N_k^i + N_{k'}^i)$ is the count-only inverse-variance choice—Aitken’s classical generalised-least-squares solution [2] restricted to count-only weights—and an unconditional Cauchy–Schwarz inequality shows it minimises the delta-method variance *surrogate* $V_k + V_{k'}$ for any impression counts. We propose the plug-in delta-method asymptotic variance $\tilde{V}(\omega)$ of the self-normalised weighted ratio estimator [11, 13] as a finite-sample diagnostic that is well aligned with the conditional-on-counts efficiency question, in contrast to a naive pair-bootstrap. On the full KDD Cup 2012 Track 2 sponsored-search log (235M impressions, 4.6M interventional pairs), the plug-in asymptotic relative efficiency of harmonic over uniform is 4.53, well-calibrated against click-resample standard error (1–2% at two sample fractions), and robust to weight-capping (ARE ≥ 4.1 at the 99.99th percentile). A Cochran- Q diagnostic confirms substantial PBM mis-specification ($Q/\text{df} = 3.83$), allowing efficiency comparisons to be interpreted separately from model fit; under such mis-specification different weight schemes target different weighted averages of cell-specific ratios—an estimand caveat we discuss explicitly. Synthetic (43%–54% MSE reduction at high imbalance) and Yahoo-LTR semi-synthetic ($\sim 40\%$ variance reduction at $r = 50$) experiments with ground truth corroborate the efficiency prediction.

ACM Reference Format:

Anonymous Author(s). 2026. Count-Only Inverse-Variance Weighting for Intervention Harvesting. In *Proceedings of ACM CIKM (CIKM '26)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Position bias systematically corrupts click signals used for learning-to-rank [8]; intervention harvesting [1] estimates position-bias ratios $R_{k,k'} = p_k/p_{k'}$ from naturally occurring ranking variations without explicit randomisation. On production click logs, the per-cell counts that feed this estimator are extraordinarily heavy-tailed. On the full KDD Cup 2012 Track 2 release (235M impressions; 4.6M interventional pairs at adjacent positions $(1, 2)$), **60.7% of pairs**

have $N_1 + N_2 \leq 5$ while 0.1% of pairs carry 42% of all impressions [9]. Public benchmarks hide this shape: the released Baidu-ULTR [16] aggregation has 94% of (q, d, pos) cells with $N = 1$ and 78% of adjacent-pair impression ratios within $[0.5, 2]$ *because* the publication pipeline deduplicates impressions per cell—an artefact of release engineering, not of the underlying production logs. The unaggregated KDD release preserves the heavy-tailed production distribution.

Equal-weighting heteroscedastic estimators is statistically sub-optimal by a century of classical theory: Aitken’s generalised-least-squares theorem [2] mandates inverse-variance weights, a result the meta-analysis literature [5, 7] has championed for fifty years. Recent ULTR work [1, 14] introduced *harmonic* weighting $\omega_i = N_k^i N_{k'}^i / (N_k^i + N_{k'}^i)$ to reduce variance versus uniform on heavy-tailed counts, but existing validation has relied on *bootstrap* resampling, which we show can be misleading for inverse-variance-style weights on singleton-dominated heavy-tailed data (Section 3).

Contributions. (1) We identify harmonic as the count-only inverse-variance choice under the standard PBM with homoscedastic relevance (Theorem 2), prove an unconditional Cauchy–Schwarz reduction of the delta-method variance *surrogate* $V_k + V_{k'}$ (Theorem 3, conditional on the surrogate objective and unconditional in the counts), and characterise the exact asymmetric surrogate oracle of which harmonic is the $\alpha = \beta$ special case (Theorem 4). (2) We adopt the plug-in delta-method asymptotic variance $\tilde{V}(\omega)$ as a finite-sample efficiency diagnostic aligned with the conditional-on-counts question, and use empirical click-resample SE as its calibration baseline; we discuss why pair-bootstrap is misaligned with this question on heavy-tailed data. (3) On the full KDD log we measure $\text{ARE}(\text{harmonic}, \text{uniform}) = 4.53$, calibrated within 1–2% across two sample fractions, and robust to weight-capping at the 99.99th percentile ($\text{ARE} \geq 4.1$). We quantify PBM mis-specification via Cochran’s $Q/\text{df} = 3.83$ and discuss the estimand caveat that follows. Synthetic and Yahoo-LTR experiments with ground truth corroborate.

2 Estimator and Theory

Under PBM, $P(\text{click} \mid q, d, k) = p_k r_{q,d}$. Given an interventional set $I_{k,k'}$ of (q, d) pairs that appeared at both positions k and k' , the position-bias ratio is estimated as [1]

$$\hat{R}_{k,k'} = \frac{A}{B} = \frac{\sum_i \omega_i \hat{\theta}_k^i}{\sum_i \omega_i \hat{\theta}_{k'}^i}, \quad \hat{\theta}_k^i = \frac{C_k^i}{N_k^i}. \quad (1)$$

This is a self-normalised weighted ratio estimator. We say a weight scheme is *ratio-unbiased* when the ratio of conditional expectations of numerator and denominator equals $p_k/p_{k'}$, i.e. $\mathbb{E}[A \mid N]/\mathbb{E}[B \mid N] = p_k/p_{k'}$ (not $\mathbb{E}[\hat{R}] = p_k/p_{k'}$; the finite-sample bias of the self-normalised ratio is $O(1/N)$ under standard regularity). Under PBM, the ratio of conditional expectations of numerator and denominator equals $p_k/p_{k'}$ for any count-only weight $\omega_i = \omega(N_k^i, N_{k'}^i)$, and

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
CIKM '26, Toronto, Canada

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

the resulting self-normalised ratio estimator is consistent for this target [1].

PROPOSITION 1 (CONDITIONAL DELTA-METHOD VARIANCE). *By the delta method [13], $\text{Var}(\hat{R} \mid \mathbf{N}) \approx R^2[\alpha V_k(\omega) + \beta V_{k'}(\omega)]$ with $V_r(\omega) = \sum_i \omega_i^2 / N_r^i / (\sum_i \omega_i)^2$, $\alpha = (1 - p_k)/p_k$, $\beta = (1 - p_{k'})/p_{k'}$.*

Harmonic as count-only IVW. For Bernoulli observations, $\text{Var}(\hat{\theta}_k^i) = \theta_k^i(1 - \theta_k^i)/N_k^i$. Under PBM, $\theta_k^i = p_k r_i$, so the Bernoulli variance contains a cell-level relevance r_i that is generally heterogeneous. After suppressing this cell-level relevance heterogeneity, the inverse-variance optimum of Aitken [2] restricted to count-only weights is the harmonic mean of the two cell counts. We make this precise as a statement about the variance *surrogate* obtained by ignoring cell-level relevance variation:

THEOREM 2 (COUNT-ONLY IVW ON THE SURROGATE). *Among all non-negative weights $\omega_i = \omega(N_k^i, N_{k'}^i)$, harmonic weighting minimises the count-only delta-method variance surrogate $V_k(\omega) + V_{k'}(\omega)$.*

THEOREM 3 (SURROGATE GUARANTEE, UNCONDITIONAL IN COUNTS). *For any impression counts $\{N_k^i, N_{k'}^i\}_{i=1}^n$, $V_k(\omega^h) + V_{k'}(\omega^h) \leq V_k(1) + V_{k'}(1)$. The guarantee is unconditional in the counts but conditional on $V_k + V_{k'}$ as the objective; it implies improved $\text{Var}(\hat{R} \mid \mathbf{N})$ under PBM when $\alpha \approx \beta$ (Proposition 1 below).*

SKETCH. Substituting $\omega_i^h = N_k^i N_{k'}^i / (N_k^i + N_{k'}^i)$ collapses the LHS to $1/W$ with $W = \sum_i \omega_i^h$. The RHS is $(1/n^2) \sum_i 1/\omega_i^h$. Cauchy-Schwarz gives $n^2 \leq W \sum_i 1/\omega_i^h$. $\square$

THEOREM 4 (EXACT ASYMMETRIC ORACLE). $\alpha V_k(\omega) + \beta V_{k'}(\omega) \geq (\sum_i 1/s_i)^{-1}$ with $s_i = \alpha/N_k^i + \beta/N_{k'}^i$, attained at $\omega_i^* \propto N_k^i N_{k'}^i / (\alpha N_{k'}^i + \beta N_k^i)$. Harmonic is the $\alpha = \beta$ special case.

Theorem 3 is unconditional in the counts – a property that distinguishes harmonic from the min-count heuristic of [1], which has no such guarantee. Theorem 4 provides the oracle benchmark; in synthetic experiments substituting the true α, β into ω^* yields variance within $\pm 2\%$ of harmonic. Complete proofs and additional sensitivity analyses are provided in the supplementary material.

Limits of the symmetric guarantee. The unconditional bound in Theorem 3 concerns $V_k + V_{k'}$. The asymmetric analogue $\alpha V_k(\omega^h) + \beta V_{k'}(\omega^h) \leq \alpha V_k(1) + \beta V_{k'}(1)$ does *not* hold unconditionally: a closed-form $n = 2$ counterexample with $(N_k^1, N_{k'}^1) = (1, 100)$, $(N_k^2, N_{k'}^2) = (100, 1)$ and $(\alpha, \beta) = (0.01, 1)$ has harmonic strictly worse than uniform on the asymmetric surrogate. The configurations under which this fails are degenerate (each cell informative at exactly one position); we have not encountered them in any realistic count distribution we tested. The ω^* in Theorem 4 remains the exact oracle for all (α, β) .

3 Plug-in Asymptotic Variance: the Right Instrument

$\tilde{V}$ in closed form. With $A = \sum_i \omega_i \theta_k^i$, $B = \sum_i \omega_i \theta_{k'}^i$, and clicks at k and k' conditionally independent given counts, the delta method gives

$$\tilde{V}(\omega) = \frac{\sum_i \omega_i^2 \theta_k^i (1 - \theta_k^i) / N_k^i}{A^2} + \frac{\sum_i \omega_i^2 \theta_{k'}^i (1 - \theta_{k'}^i) / N_{k'}^i}{B^2}, \quad (2)$$

Table 1: Plug-in SE vs. click-resample SE across KDD subsamples, and ARE(harmonic, uniform) at each scale. Click resample only computed at small f for runtime; agreement is within 1–2% where measured. ARE grows with data volume because heavy-tail mass becomes more pronounced at larger scales.

f	pairs	$\widehat{\text{SE}}_U^{\text{plug}}$	$\widehat{\text{SE}}_U^{\text{boot}}$	$\widehat{\text{SE}}_H^{\text{plug}}$	ARE
10^{-3}	4.3K	0.131	0.132	0.111	1.66
10^{-2}	54K	0.039	0.038	0.028	2.21
10^{-1}	542K	0.012	–	0.007	3.07
1.0	4.61M	0.0042	–	0.0020	4.53

$\text{SE}(\hat{R}) = |\hat{R}| \sqrt{\tilde{V}(\omega)}$. Plugging in $\hat{\theta}_i$ (with light shrinkage toward the marginal) yields the textbook self-normalised IS variance estimator [10–12].

ARE. For schemes a, b , $\text{ARE}(a, b) = \tilde{V}(\omega_b) / \tilde{V}(\omega_a)$: scheme b needs $\text{ARE}(a, b)$ times the data of a for equal precision. $\tilde{V}$ scales as $1/\text{ESS}(\omega)$, with Kong’s effective sample size [10] $\text{ESS}(\omega) = (\sum_i \omega_i)^2 / \sum_i \omega_i^2$.

Why pair-bootstrap is misaligned here. A naive bootstrap that resamples (q, d) pairs estimates a different quantity than the conditional on-counts efficiency we care about, for three reasons. (i) Averaging over $\sim 60\%$ near-singleton pairs under uniform produces a smooth resampling distribution whose tightness reflects averaging-of-noise rather than estimator precision. (ii) Pair resampling perturbs the high-information cells that drive any inverse-variance estimator, so concentrated weight schemes are systematically penalised by the resampling design, not by their statistical behaviour. (iii) Pair-bootstrap entangles estimator efficiency with model fit; $\tilde{V}$ separates them (Section 4). The appropriate empirical baseline against which to validate the plug-in is the click-resample (parametric Bernoulli bootstrap conditional on counts), which we adopt for calibration.

Calibration and stability at scale. Table 1 reports plug-in SE versus empirical click-resample SE across KDD sample fractions, together with ARE(harmonic, uniform) at each scale. Plug-in agrees with empirical click-resample SE within 1–2% where the latter is computed, and the ARE grows with sample size ($1.66 \rightarrow 4.53$), because the heavy-tailed structure of the count distribution becomes more pronounced as more data is included rather than averaged out.

4 Real-Log Results: KDD Cup 2012 Track 2

Data. Full release: 149.6M raw rows expanding to 235.6M impressions, 63.8M unique $(q, \text{ad}, \text{pos})$ cells, 8.2M clicks (CTR 3.49%). Pair (1, 2) interventional set: 4.61M (q, ad) pairs, 132.3M impressions, median pair size 4, max 2.47M; 60.7% of pairs have $N_1 + N_2 \leq 5$.

Headline (Table 2). $\text{ARE}(\text{harmonic, uniform}) = 4.53$. **Uniform requires $\sim 4.5\times$ the impressions of harmonic for matching precision**, with $\text{SE}_U / \text{SE}_H \approx 2.13$. The min heuristic [1] achieves 4.41, the count-only oracle 4.57; harmonic is the principled parameter-free optimum.

Table 2: Plug-in delta-method asymptotic variance $\tilde{V}$, SE of $\hat{R}$, ARE vs. uniform, and ESS on full KDD Cup 2012 Track 2 ($n_{\text{pairs}} = 4.61\text{M}$). Harmonic, min, and the count-only oracle all reach ARE ≈ 4.5 . The full- θ oracle is pathological because per-cell $\hat{\theta}_i$ is too noisy at this scale.

Scheme	$\hat{R}$	SE	ARE	ESS
Uniform	1.683	4.24e-3	1.00	4.61M
Min [1]	1.712	2.05e-3	4.41	2,526
Harmonic	1.710	2.02e-3	4.53	1,849
Count-oracle (0.5,1)	1.715	2.02e-3	4.57	1,760
θ -aware oracle	2.343	0.148	0.002	1.1

Table 3: Stratified ARE(harmonic, uniform) by $N_1 + N_2$ tertile on full KDD (the bottom quartile boundary collapses into the lowest tertile because singletons dominate). Schemes are tied on the noise-dominated strata; harmonic dominates $\sim 5\times$ on the high-information stratum where the actual estimation signal lives.

$N_1 + N_2$ tertile	n pairs	$\tilde{V}_U$	ARE(h, u)
Lowest	1.99M	2.3e-5	1.02
Middle-low	1.38M	1.9e-5	1.09
Highest	1.25M	7.6e-6	4.98

Stratified ARE (Table 3). Schemes tie on noise singletons (where there is nothing to estimate) and harmonic dominates uniform by $5\times$ on the high-information stratum. The gain comes from *not* treating singletons as equal to million-impression cells.

Weight-cap robustness. The harmonic ESS of 1,849 out of 4.61M pairs invites a concern that the efficiency gain rests on a few extreme cells, but the small ESS reflects deliberate concentration on the pairs containing nearly all statistical information in this heavy-tailed setting rather than instability of the estimator. Capping ω_i^h at the 99.99th percentile (threshold 2,079; 462 pairs capped) lifts ESS to 5.2×10^4 and still gives ARE = 4.12; at the 99th percentile, ARE = 2.83 (ESS 7.6×10^5); even at the 90th, ARE = 1.86. Harmonic beats uniform robustly across capping levels.

Comparison of count-only weighting schemes. Table 4 compares harmonic to the natural alternatives on full KDD. Among count-only weights, harmonic is best on $\tilde{V}$. Log-counting is too flat (treats singletons nearly as well as million-impression cells); the geometric mean and min-count concentrate more aggressively but neither matches the inverse-variance optimum. Capping at the 99th percentile preserves a $2.83\times$ advantage but shifts the estimand ($\hat{R} : 1.71 \rightarrow 1.61$), consistent with the estimand discussion below.

Estimand under PBM mis-specification. PBM is a known approximation; on real logs different weight schemes estimate different functionals of cell-specific ratios. Specifically, the population limit of $\hat{R}$ under any weight ω is

$$R_\omega = \frac{\sum_i \omega_i \theta_k^i}{\sum_i \omega_i \theta_{k'}^i},$$

Table 4: Count-only weighting schemes compared on full KDD via plug-in $\tilde{V}$. Schemes that concentrate more aggressively give lower $\tilde{V}$ but smaller ESS. Harmonic is the inverse-variance optimum among count-only schemes (Theorem 2).

Scheme	$\hat{R}$	SE	ARE	ESS
Uniform	1.683	4.24e-3	1.00	4.61M
$\log(1 + \min(N_k, N_{k'}))$	1.651	3.18e-3	1.71	3.17M
$\sqrt{N_k N_{k'}}$	1.706	2.28e-3	3.55	1,246
$\min(N_k, N_{k'})$	1.712	2.05e-3	4.41	2,526
Harmonic	1.710	2.02e-3	4.53	1,849
Harmonic, cap@99th pct	1.606	2.40e-3	2.83	765K

which equals $p_k/p_{k'}$ only when $\theta_k^i/\theta_{k'}^i$ is the same constant across cells—i.e. when PBM holds exactly. Under mis-specification each scheme thus targets a different weighted average of the cell-specific ratios $\rho_i = \theta_k^i/\theta_{k'}^i$ (uniform $\rightarrow 1.68$, harmonic $\rightarrow 1.71$, θ -aware oracle $\rightarrow 2.34$ on KDD). The plug-in $\tilde{V}$ comparison is a precision statement at fixed ω —a model-based efficiency comparison—and does *not* claim that one R_ω is the “correct” position-bias ratio. Choosing ω thus involves a precision/estimand trade-off; harmonic’s attractive property is that, among count-only schemes, it concentrates on high-information cells (where ρ_i is estimated most precisely) without depending on noisy plug-in $\hat{\theta}_i$.

Quantifying PBM mis-specification (Cochran’s Q). PBM is a known approximation, not a literal model – prior work [6, 14, 15] has long established that real logs violate it. The relevant question is how large the deviation is. Cochran’s Q on the 16,132 high-information pairs (≥ 5 clicks at each of positions 1 and 2) gives $Q = 6.17\times 10^4$, $\text{df} = 16,131$, $Q/\text{df} = 3.83$: per-cell odds-ratio variability is $\sim 3.8\times$ what binomial sampling alone would produce, well outside Poisson dispersion. This magnitude is the quantitative driver of the estimand drift discussed in Paragraph 4; plug-in $\tilde{V}$ keeps the efficiency comparison well-posed by working at fixed ω .

5 Synthetic and Yahoo-LTR Validation

To validate the ARE prediction in settings with ground truth, we run two controlled experiments.

Anchor-item imbalance sweep (synthetic). Following [1]: $M=10$ positions, $p_k = k^{-\eta}$ with $\eta = 1$, AdjacentChain estimator. A fraction s of documents are “anchors” with heavy traffic at one position ($N \sim \text{Unif}(50r, 200r)$) and light traffic at the other ($N \sim \text{Unif}(1, 51)$); the remaining $1-s$ are balanced ($N \sim \text{Unif}(25, 75)$ at both positions); $s \in \{0.5, \dots, 0.95\}$, $r \in \{1, \dots, 100\}$, 500 trials per cell. At $s = 0.80$, $r = 100$ harmonic reduces variance by 43.6% vs. uniform and 5% vs. min; reduction grows to $\sim 54\%$ at $s = 0.90$, $r = 50$. Across the full (s, r) grid the data-equivalence factor $\text{DEF}(s, r) = b_{\text{uni}}(s, r)/b_{\text{harm}}(s, r)$ – the ratio of base impressions per document at which each scheme first hits a fixed MSE threshold – has median 1.31, mean 1.30, max 1.94. Highly imbalanced cells achieve up to $2\times$ data savings; near-balanced cells ($r = 1$) saturate at $\text{DEF} \approx 1$.

Bias comparability under PBM violations. We additionally simulate three explicit PBM violation models – trust bias, cascade, and position-dependent relevance [1?] – at imbalance ratio 50:1, split

Table 5: MSE decomposition under three explicit PBM violations (anchor imbalance 50:1, split 80/20, 5,000 trials). Count-only schemes share comparable bias² within each model; harmonic achieves the lowest variance, hence the lowest MSE.

Model	Scheme	MSE	Variance	Bias ²
Trust	Uniform	0.00487	0.00486	1.0e-5
	Min	0.00298	0.00297	2.0e-5
	Harmonic	0.00272	0.00272	0.0
Cascade	Uniform	0.01376	0.00276	0.01100
	Min	0.01242	0.00171	0.01071
	Harmonic	0.01252	0.00147	0.01105
Pos.-dep.	Uniform	0.04086	0.03973	0.00113
	Min	0.03896	0.03770	0.00127
	Harmonic	0.03724	0.03613	0.00111

80/20, 5,000 trials. Across all three models the count-only schemes (uniform, min, harmonic) exhibit *comparable* bias², so harmonic’s lower variance translates directly into lower MSE (Table 5). Under cascade, where history-dependent examination introduces a substantial structural mismatch (bias² $\approx$ 0.011 for every ratio-unbiased scheme), the variance edge is preserved but the practical MSE gain shrinks because bias dominates.

Yahoo-LTR semi-synthetic. Same simulator with graded relevance from the Yahoo Learning to Rank Challenge Set 1 test split (6,936 queries; relevance grades $g \in \{0, \dots, 4\}$ mapped to $r(q, d) = (2^g - 1)/15$ following NDCG; click model $P(C = 1 \mid q, d, k) = p_k r(q, d)$; 200 trials). Harmonic reduces variance by $\sim 40\%$ vs. uniform at imbalance $r = 50$ (Var $1.5 \times 10^{-4} \rightarrow 8.7 \times 10^{-5}$). Aggressive clipping trades variance for bias and yields $\sim 458\times$ higher MSE, confirming that variance-reduction without bias control is not viable.

Production-style traffic simulation. $K=5$ concurrent rankers; a default ranker shows each document at its natural position, treatment rankers swap adjacent pairs with probability 0.5. Under uniform traffic ($5 \times 20\%$) only 0.4% of adjacent pairs exceed impression ratio 10 and harmonic’s MSE reduction is 6.3%; under production traffic ($80\% + 4 \times 5\%$), 16.7% of pairs exceed ratio 10 and harmonic’s MSE reduction grows to **17.2%**, confirming the gain scales directly with traffic-induced imbalance.

6 Related Work

Intervention harvesting was introduced by [1] with the uniform and min-count variants we compare against; [14] jointly estimates relevance and propensities via EM. Self-normalised weighted ratio estimators have a long history in importance sampling and survey statistics: Owen [11] §9.2 derives the asymptotic variance we use; Kong [10] introduces ESS; Hesterberg [4] analyses heterogeneous-weight variance. Inverse-variance weighting traces to Aitken [2] and is canonical in meta-analysis [5, 7]. The closest analogue in counterfactual learning is the self-normalised IPS estimator SNIPS [12], which applies the same $\tilde{V}$ framework to off-policy evaluation in bandits; [3] analyses why weight heterogeneity, not sample count,

governs effective precision on logged feedback. PBM violations in real logs are documented in [14, 15] and reviewed in [6].

7 Conclusion

Framing harmonic weighting as the count-only inverse-variance optimum (Theorems 2–4) and replacing bootstrap with the plug-in delta-method variance $\tilde{V}$ (Eq. 2) yields rigorous, calibration-verified efficiency comparisons on heavy-tailed real click logs. On the full KDD Cup 2012 Track 2 log, ARE(harmonic, uniform) = 4.53, rising to $\sim 5\times$ on information-bearing cells, with PBM misspecification quantified separately by Cochran- Q ($Q/df = 3.83$). Future work: shrinkage variants of the θ -aware oracle (Theorem 4 with $\hat{\theta}_i$ regularised); click models beyond PBM [6, 14, 15] where the intervention-harvesting family must be reformulated.

GenAI Usage Disclosure

The authors used a large-language-model assistant for code generation (KDD aggregation, plug-in variance script, LaTeX table formatting), and for copy-editing. All experimental design, mathematical derivations (Theorems 2–4 and Eq. (2)), interpretation, citations, and final wording were authored and verified by the authors. No LLM-generated content was used without human review. Full code, scripts, and data-processing pipelines are provided as supplementary material.

References

- [1] A. Agarwal, X. Wang, C. Li, M. Bendersky, and M. Najork. Estimating position bias without intrusive interventions. In *WSDM*, 2019.
- [2] A. C. Aitken. On least squares and linear combination of observations. *Proc. Royal Soc. Edinburgh*, 55:42–48, 1934.
- [3] L. Bottou, J. Peters, J. Quiñero-Candela, et al. Counterfactual reasoning and learning systems. *JMLR*, 14:3207–3260, 2013.
- [4] T. Hesterberg. Weighted average importance sampling and defensive mixture distributions. *Technometrics*, 37(2):185–194, 1995.
- [5] M. Borenstein, L. V. Hedges, J. P. T. Higgins, and H. R. Rothstein. *Introduction to Meta-Analysis*. Wiley, 2009.
- [6] A. Chuklin, I. Markov, and M. de Rijke. *Click Models for Web Search*. Morgan & Claypool, 2015.
- [7] L. V. Hedges and I. Olkin. *Statistical Methods for Meta-Analysis*. Academic Press, 1985.
- [8] T. Joachims, A. Swaminathan, and T. Schnabel. Unbiased learning-to-rank with biased feedback. In *WSDM*, 2017.
- [9] Tencent. KDD Cup 2012 Track 2: predict the click-through rate of ads. Competition release, 2012.
- [10] A. Kong. A note on importance sampling using standardized weights. Tech. Rep. 348, U. Chicago Statistics, 1992.
- [11] A. B. Owen. *Monte Carlo Theory, Methods and Examples*, ch. 9. 2013.
- [12] A. Swaminathan and T. Joachims. The self-normalized estimator for counterfactual learning. In *NeurIPS*, 2015.
- [13] A. W. van der Vaart. *Asymptotic Statistics*, ch. 3. Cambridge, 1998.
- [14] X. Wang, M. Bendersky, D. Metzler, and M. Najork. Position bias estimation for unbiased learning to rank in personal search. In *WSDM*, 2018.
- [15] Y. Yue, R. Patel, and H. Roehrig. Beyond position bias: examining result attractiveness as a source of presentation bias. In *WWW*, 2010.
- [16] L. Zou et al. A large-scale search dataset for unbiased learning to rank. In *NeurIPS*, 2022.